



Databases on AWS

Chetan Agrawal, Solutions Architect

Deven Suri, Account Manager



Agenda

- AWS Database Services
- Traditional vs AWS Data services model
- Amazon RDS
- Amazon Aurora
- Amazon DynamoDB
- Amazon ElastiCache
- Amazon Neptune
- Amazon DocumentDB
- Amazon QLDB
- Amazon Timestream

Common data categories and use cases



Relational

Referential integrity, ACID transactions, schema-on-write

Lift and shift, ERP, CRM, finance



Key-value

High throughput, low-latency reads and writes, endless scale

Real-time bidding, shopping cart, social, product catalog, customer preferences



Document

Store JSON documents with quick access, query on any attribute

Content management, personalization, mobile



In-memory

Query by key with microsecond latency

Leaderboards, real-time analytics, caching



Graph

Quickly and easily create and navigate relationships between data

Fraud detection, social networking, recommendation engine



Time-series

Collect, store, and process data sequenced by time

IoT applications, event tracking



Ledger

Complete, immutable, and verifiable history of all changes to application data

Systems of record, supply chain, health care, registrations, financial



Warehouse

High performance querying on large volumes of data

Analytics, Data Marts

Common data categories and use cases



Relational

Referential integrity, ACID transactions, schema-on-write



Key-value

High throughput, low-latency reads and writes, endless scale



Document

Store JSON documents with quick access, query on any attribute



In-memory

Query by key with microsecond latency



Graph

Quickly and easily create and navigate relationships between data



Time-series

Collect, store, and process data sequenced by time



Ledger

Complete, immutable, and verifiable history of all changes to application data



Warehouse

High performance querying on large volumes of data



RDS



DynamoDB



DocumentDB



ElastiCache

Redis Memcached



Neptune



Timestream

Preview




QLDB



Redshift

Traditional Database Architecture


*one
database
for all
workloads*



Traditional Database Architecture

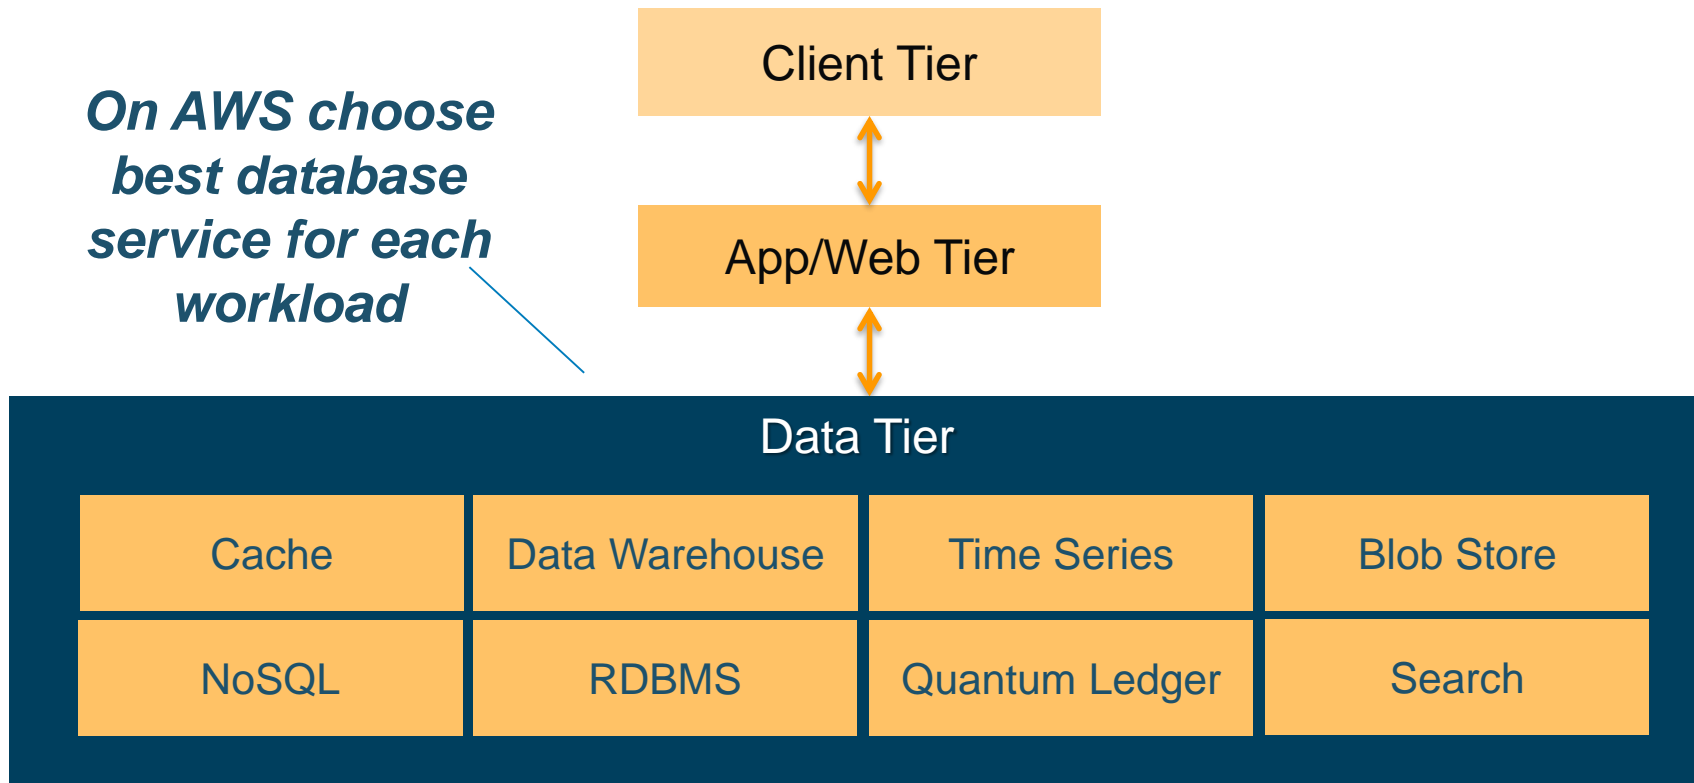
Key-value access
Complex queries
OLAP transactions
Analytics

*All forced into the
relational database*

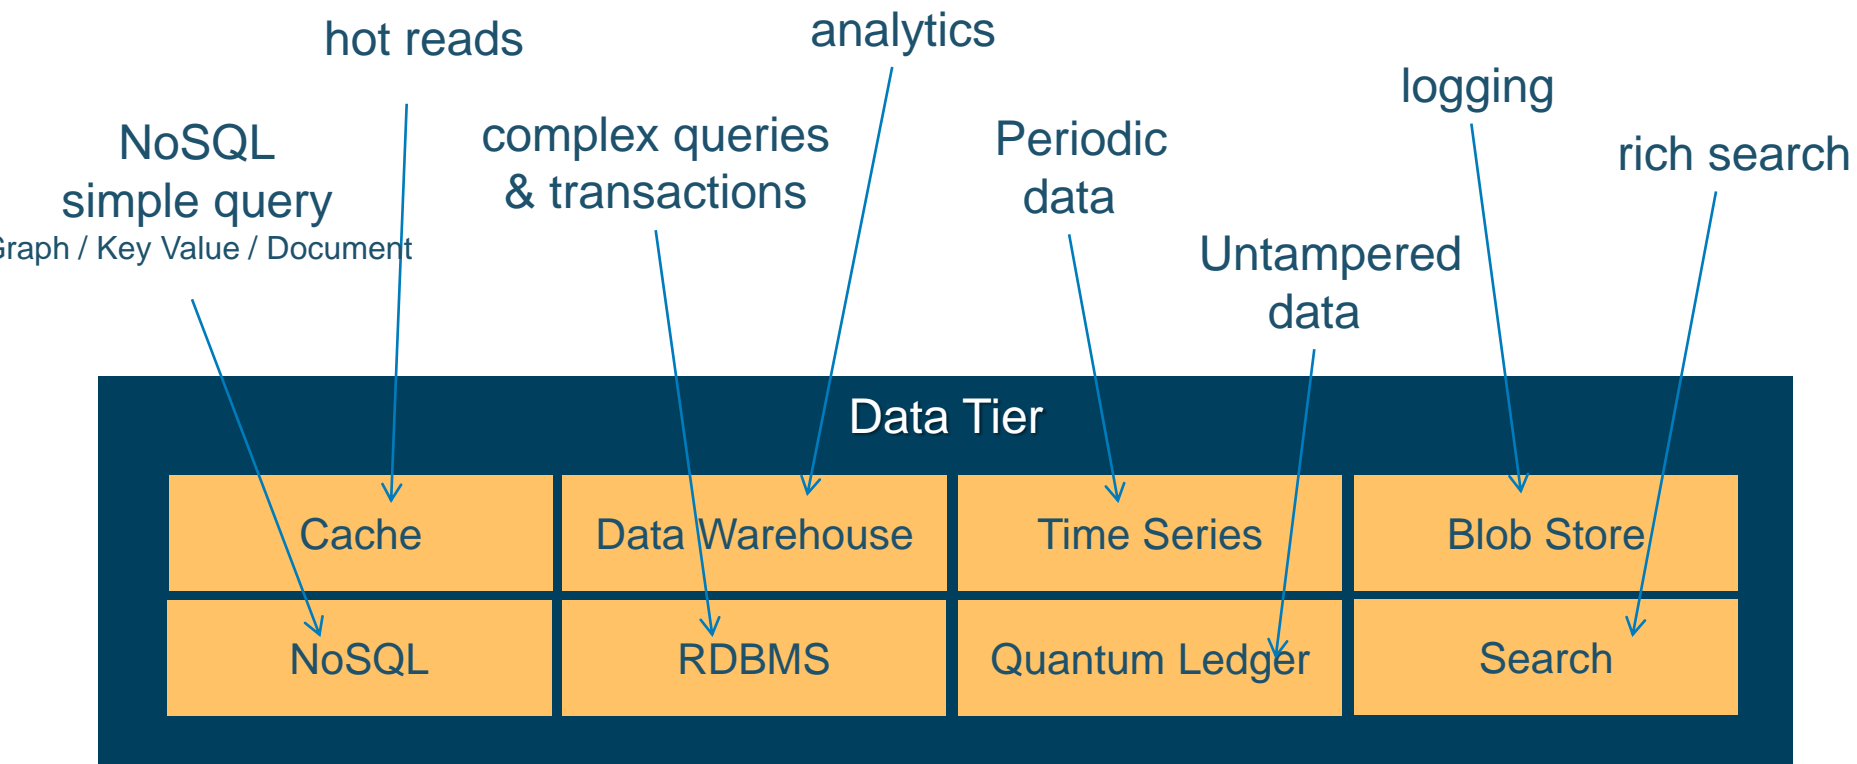


AWS Data Tier Architecture

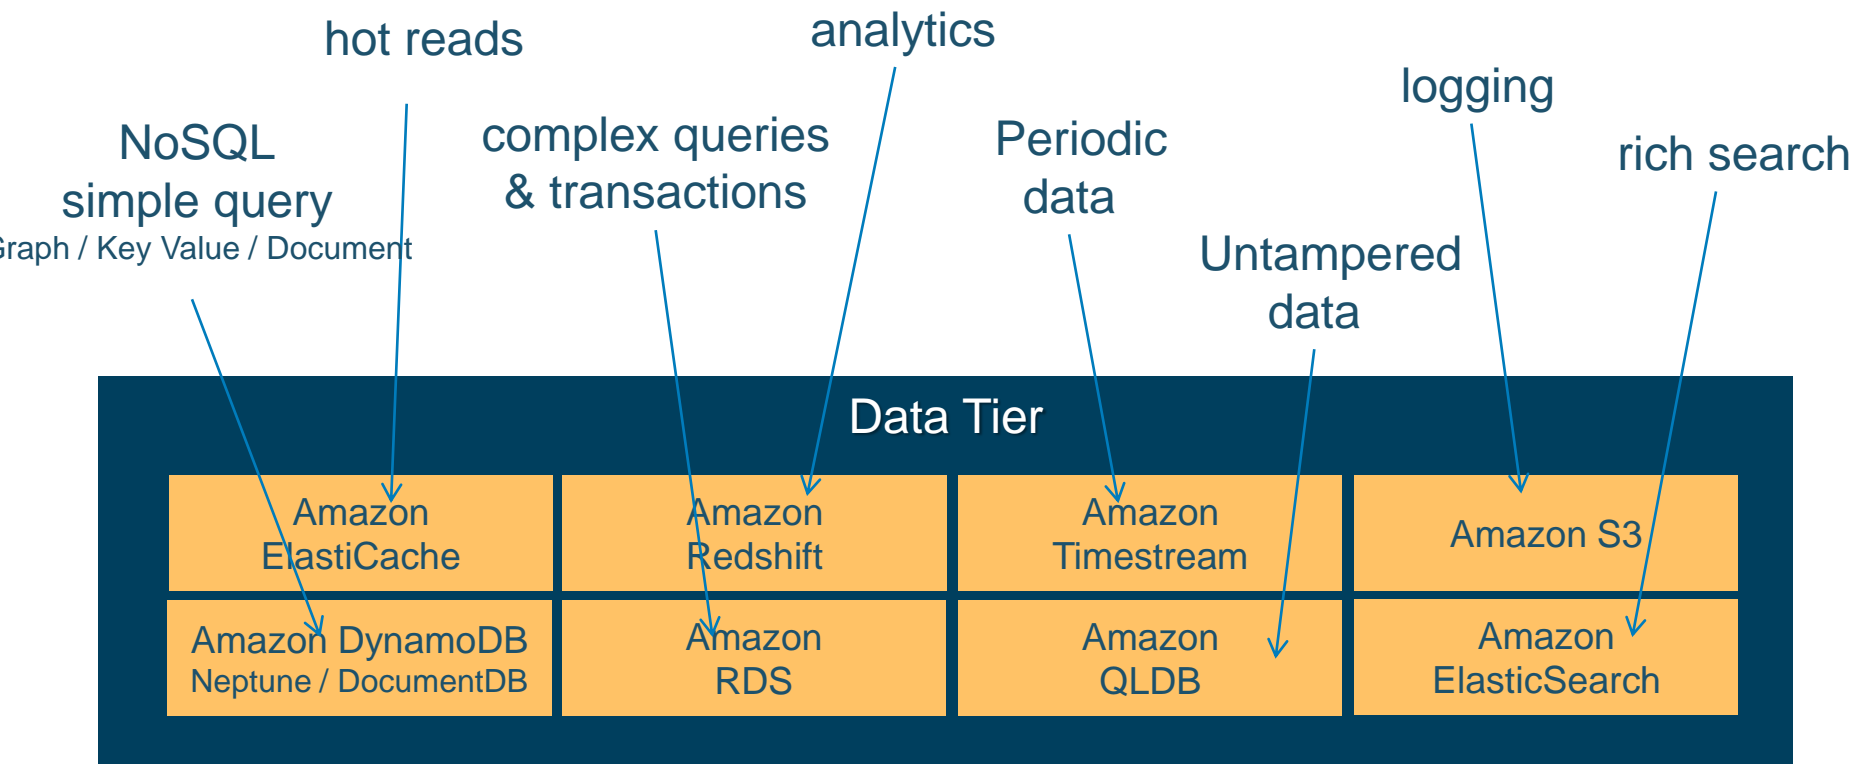
*On AWS choose
best database
service for each
workload*



Workload Driven Data Store Selection



AWS Database Services for the Data Tier



Amazon RDS

Managed relational database service with a choice of popular database engines

Amazon
Aurora



Microsoft SQL Server

ORACLE



Easy to administer

Easily deploy and maintain hardware, OS and DB software; built-in monitoring



Performant & scalable

Scale compute and storage with a few clicks; minimal downtime for your application



Available & durable

Automatic Multi-AZ data replication; automated backup, snapshots, and failover



Secure and compliant

Data encryption at rest and in transit; industry compliance and assurance programs



If you host your databases on-premises...

App optimization

Scaling

High availability

Database backups

DB s/w patches

DB s/w installs

OS patches

OS installation

Server maintenance

Rack & stack

Power, HVAC, net



you



If you host your databases in Amazon EC2...

- App optimization
- Scaling
- High availability
- Database backups
- DB s/w patches
- DB s/w installs
- OS patches
- OS installation
- Server maintenance
- Rack & stack
- Power, HVAC, net



you

- OS installation
- Server maintenance
- Rack & stack
- Power, HVAC, net



If you choose Amazon RDS...

App optimization

Scaling

High availability

Database backups

DB s/w patches

DB s/w installs

OS patches

OS installation

Server maintenance

Rack & stack

Power, HVAC, net



you

Scaling

High availability

Database backups

DB s/w patches

DB s/w installs

OS patches

OS installation

Server maintenance


Rack & stack

Power, HVAC, net




Key Amazon RDS Features


Amazon RDS Configuration	Improve Availability	Increase Throughput	Reduce Latency
Push-Button Scaling		✓	
Multi AZ	✓		
Read Replicas		✓	
Provisioned IOPS		✓	✓




Push-Button Scaling



Read Replicas

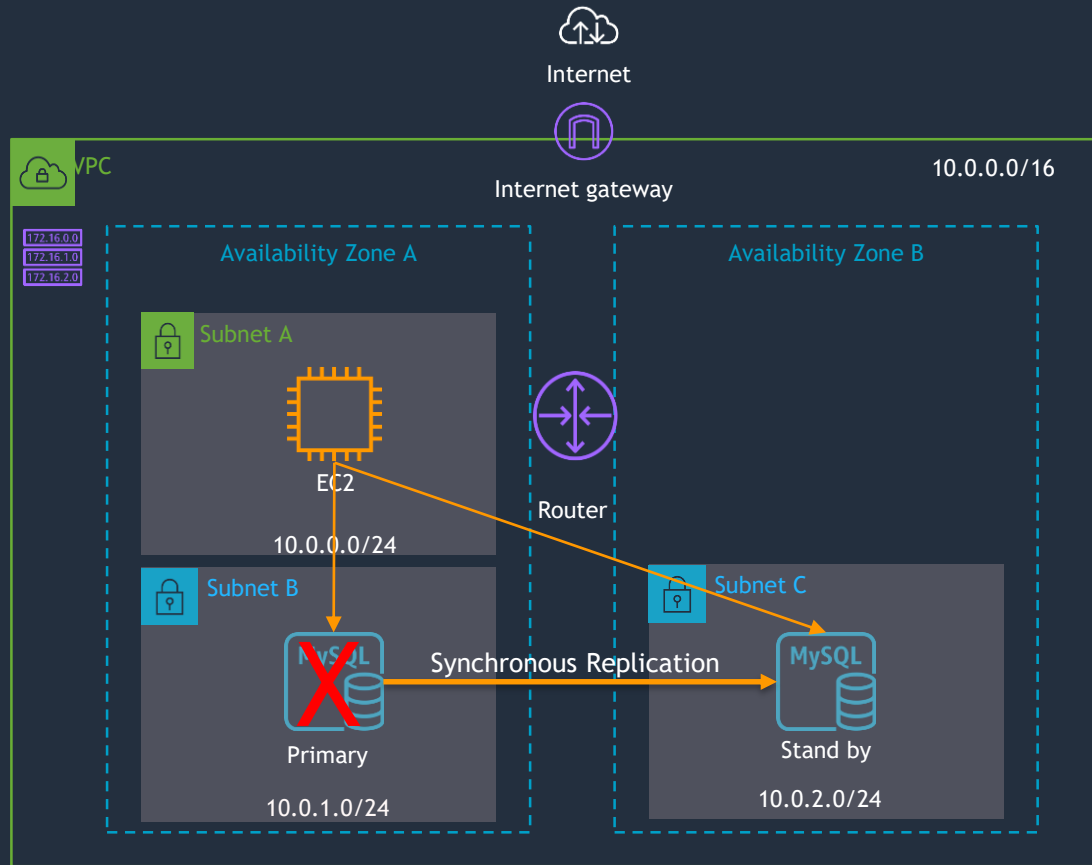


Provisioned IOPS



Demo – Multi-AZ RDS

- Launch MySQL DB in multiple Azs
- Connect to Primary
- Simulator the failure
- Check if secondary DB takes over automatically

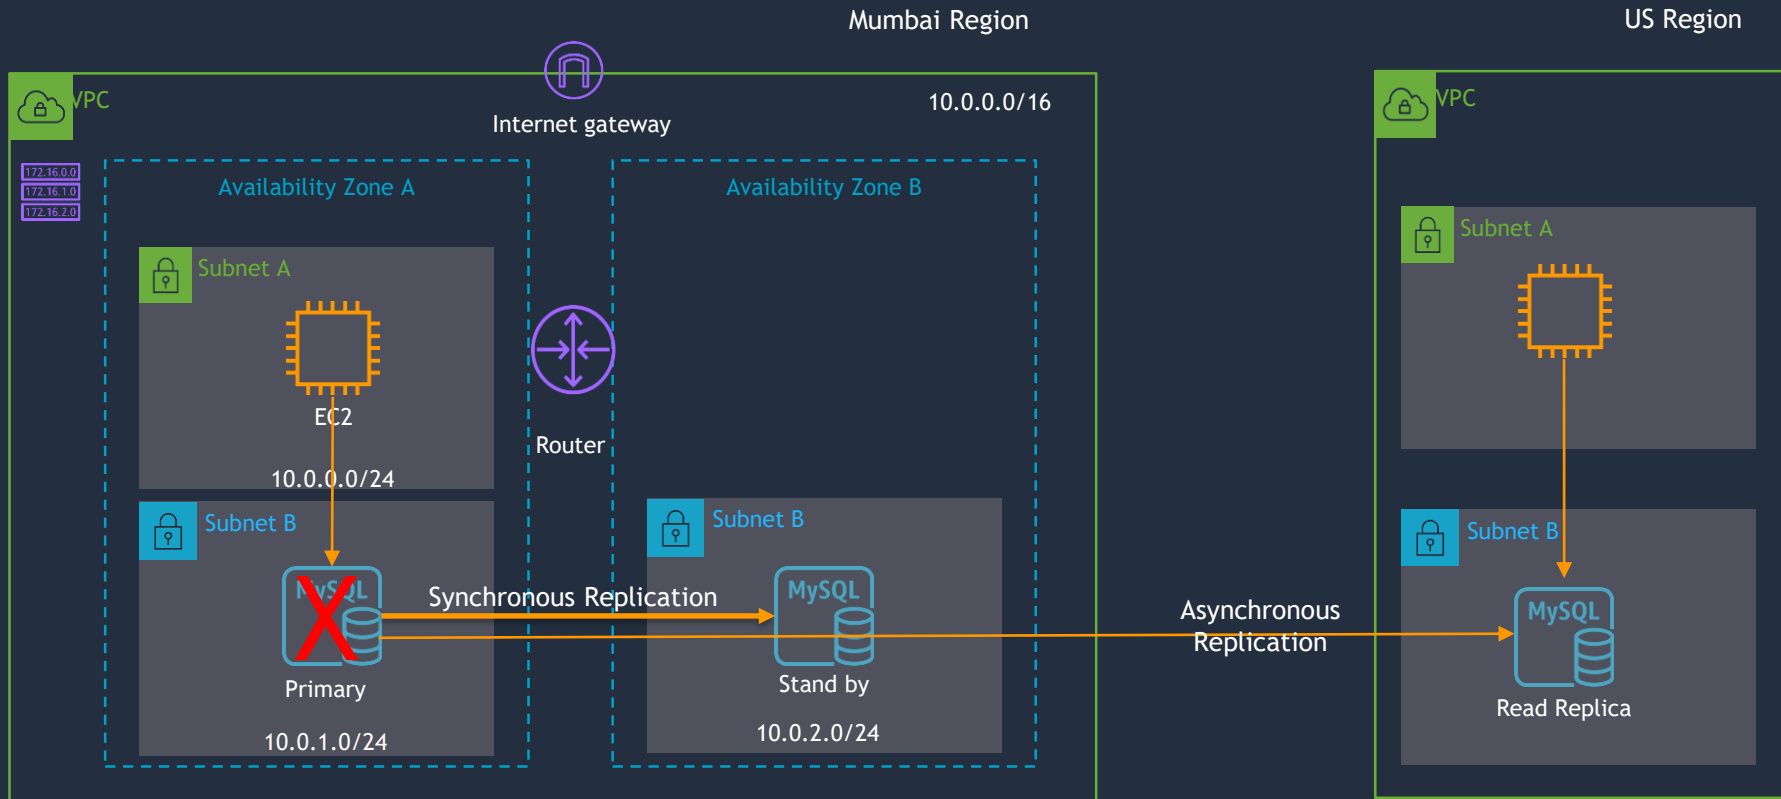


Demo steps

1. Connect to DB
2. Create table and insert some data
 - > create database test
 - > use test;
 - > create table amazon (name VARCHAR(30), id INT(2), phone VARCHAR(10), email VARCHAR(30));
 - > insert into amazon values ('Chetan Agrawal', 3, '9x2x5x3x6x', 'agrcheta@amazon.com');
3. Simulate the failure in master (Reboot with failover)
4. Wait and check if failover happens automatically

```
$while true; do host database-1.xxxxxx.ap-south-1.rds.amazonaws.com; sleep 3; done
```

Exercise – Read Replica across Region



Amazon Aurora

MySQL and PostgreSQL compatible relational database built for the cloud

Performance and availability of commercial-grade databases at 1/10th the cost




Performance & scalability

5x throughput of standard MySQL and 3x of standard PostgreSQL; scale-out up to 15 read replicas




Availability & durability

Fault-tolerant, self-healing storage; six copies of data across three AZs; continuous backup to S3



Highly secure

Network isolation, encryption at rest/transit

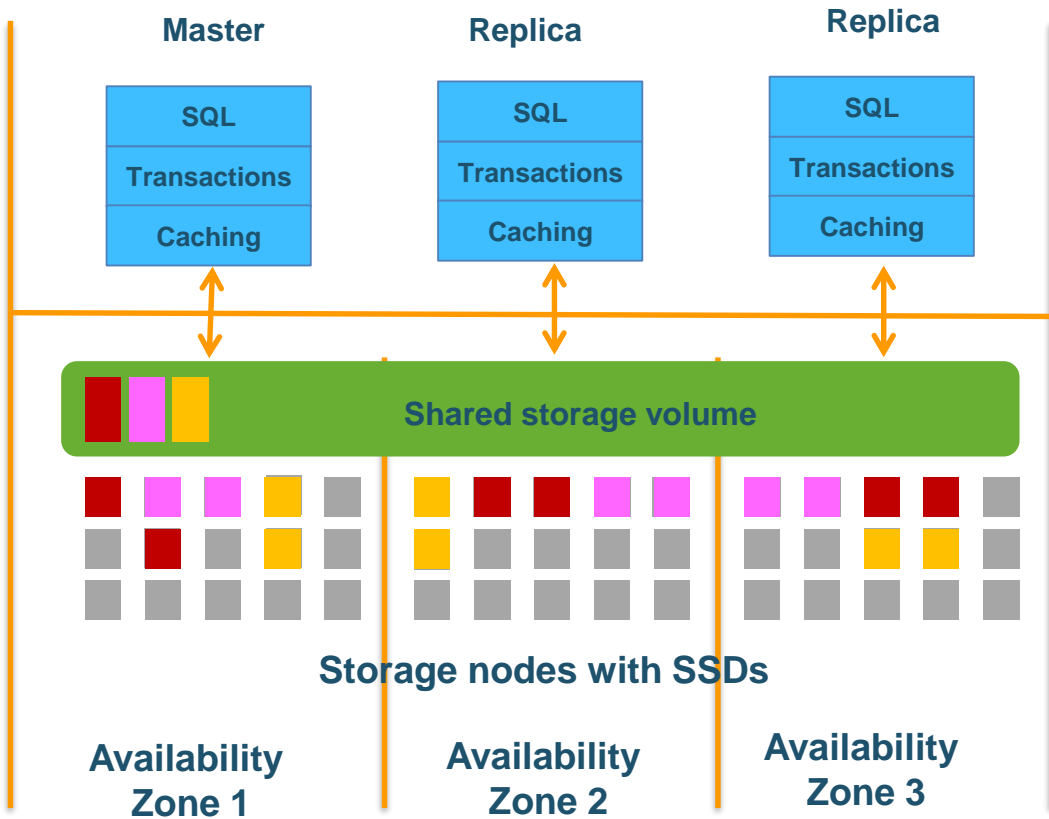


Fully managed

Managed by RDS: no hardware provisioning, software patching, setup, configuration, or backups


Scale-out, distributed, multi-tenant architecture

- Purpose-built log-structured distributed storage system designed for databases
- Storage volume is striped across hundreds of storage nodes distributed over 3 different Availability Zones
- Six copies of data, two copies in each Availability Zone to protect against AZ+1 failures
- Master and replicas all point to the same storage




Aurora MySQL performance

WRITE PERFORMANCE



READ PERFORMANCE



MySQL SysBench results; R4.16XL: 64cores / 488 GB RAM

Aurora



MySQL 5.6

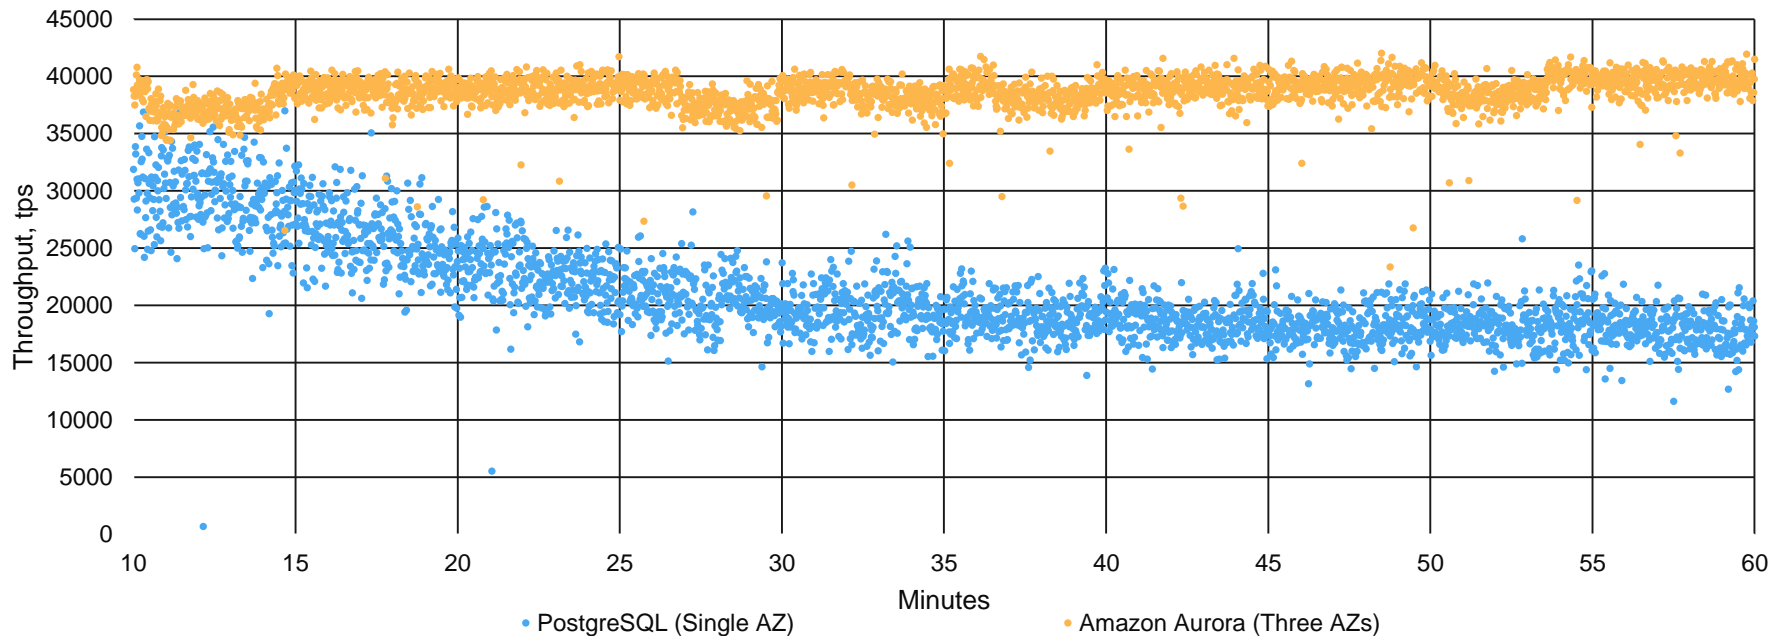


**Aurora read write throughput compared to MySQL 5.6
based on industry standard benchmarks.**

Aurora PostgreSQL performance

While running pgbench at load, throughput is 3x more consistent than PostgreSQL

pgbench throughput over time, 150 GiB, 1024 clients




...and more



Up to 64TB of storage – auto-incremented in 10GB units

- Automatic storage scaling up to 64 TB—no performance impact
- Continuous, incremental backups to Amazon S3
- Instantly create user snapshots—no performance impact
- Automatic restriping, mirror repair, hot spot management, encryption

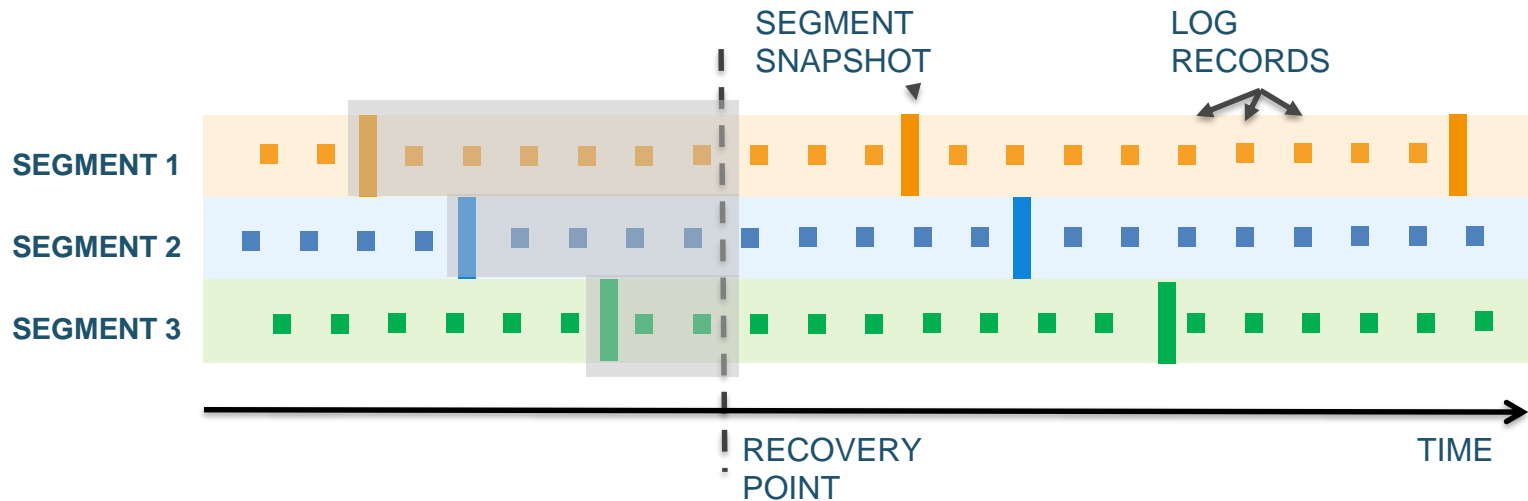
Database backtrack



Backtrack brings the database to a point in time without requiring restore from backups

- Backtracking from an unintentional DML or DDL operation
- Backtrack is not destructive. You can backtrack multiple times to find the right point in time

How does backtrack work?



We keep periodic snapshot of each segment; we also preserve the redo logs

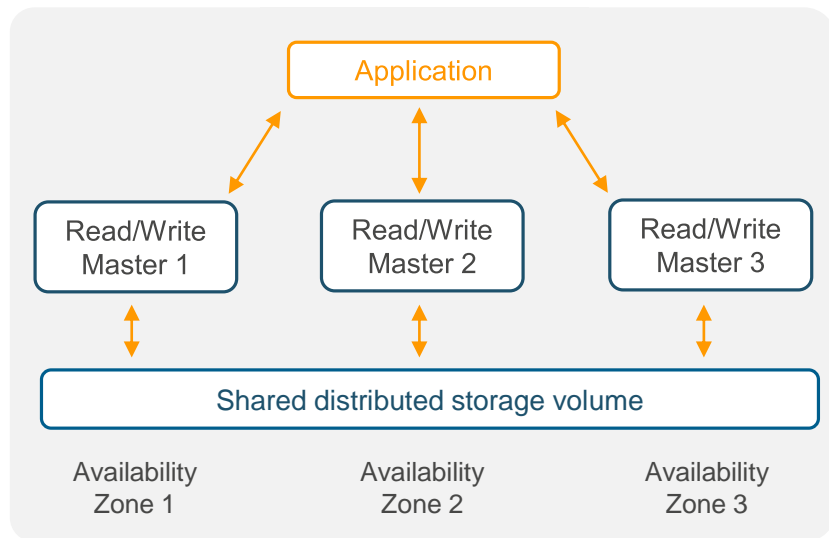
For backtrack, we identify the appropriate segment snapshots

Apply log streams to segment snapshots in parallel and asynchronously

Aurora Multi-Master

First relational database service with scale-out reads and writes across multiple data centers

Scale out both reads **and** writes



Zero application downtime from ANY instance failure

Zero application downtime from ANY AZ failure

Faster write performance and higher scale

Global database

Faster disaster recovery and enhanced data locality

Promote read-replica to a master for faster recovery **in the event of disaster**

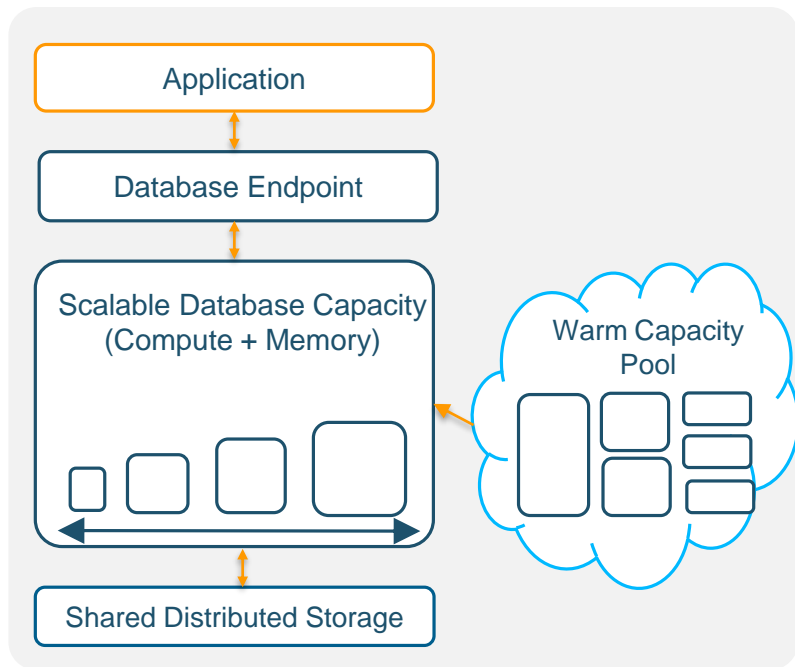
Bring data close to your customer's applications in **different regions**

Promote to a master for **easy migration**



Aurora Serverless

On-demand, auto-scaling database for applications with variable workloads



Starts up on demand, shuts down when not in use

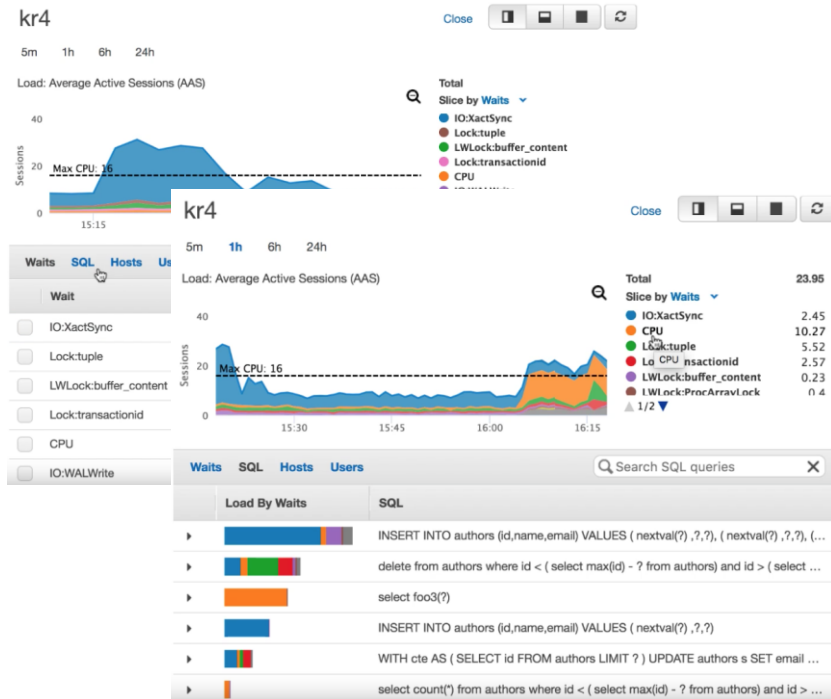
Automatically scales with no instances to manage


Pay per second for the database capacity you use

Performance Insights for Aurora

Analyze and troubleshoot your database performance

- Supports PostgreSQL and MySQL
- Expands on existing Amazon RDS monitoring features to analyze issues and performance
- Easy bottleneck identification – keep track of performance metrics such as high CPU consumption, lock waits, I/O latency, and SQL statements





NoSQL database

Seamless scalability

Zero admin

Single-digit millisecond latency

Multi-Master

Multi-Region


Amazon DynamoDB



Fully managed



Consistently fast at any scale



Highly available
and durable



Secure

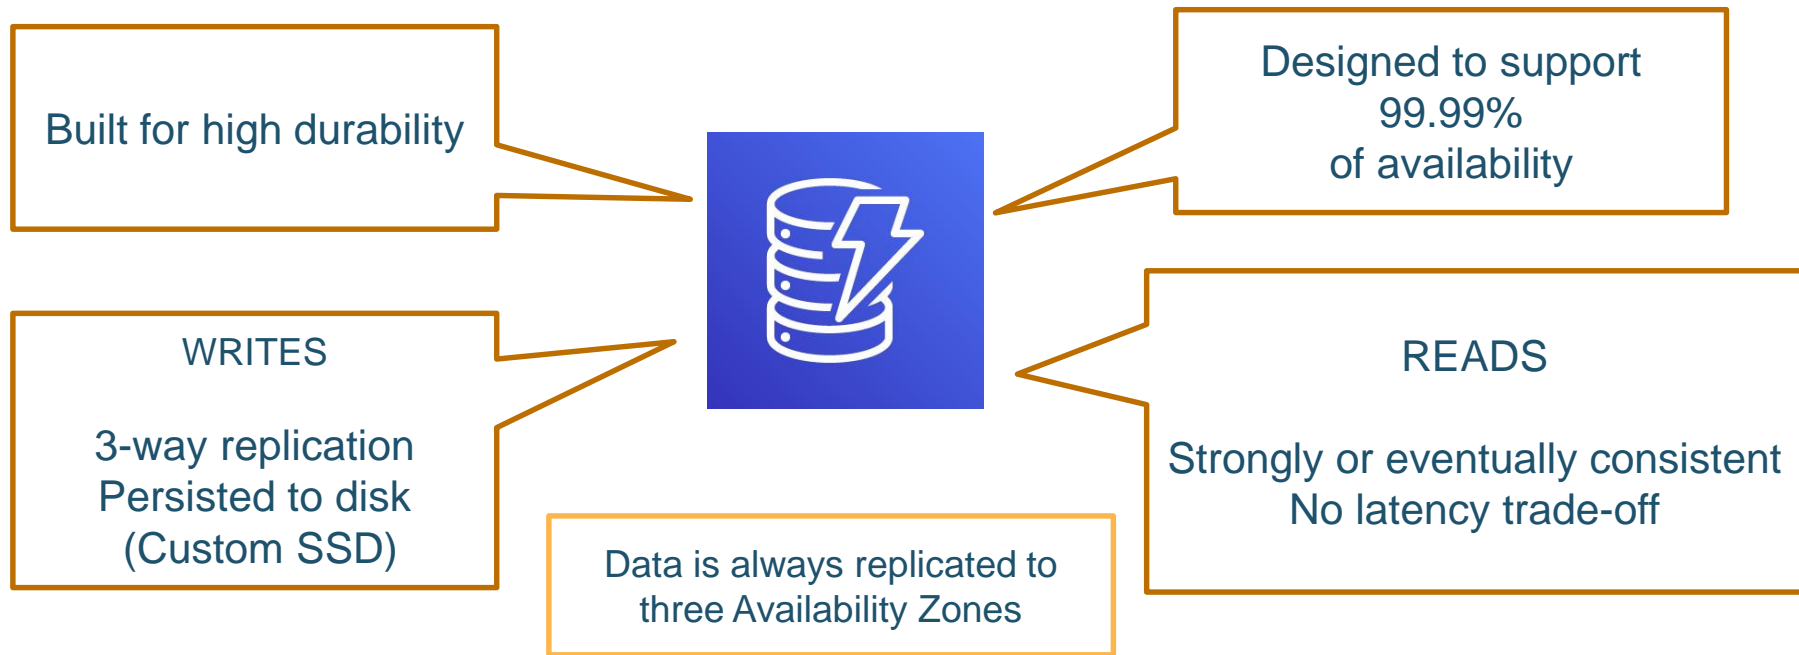


Integrates with AWS Lambda,
Amazon Redshift, and more



Cost-effective

Highly available and durable



Highly available and durable

OrderId: 1
CustomerId: 1
ASIN: [B00X4WHP5E]

3-way replication

Data is always replicated to three Availability Zones

Hash(1) = 7B



Availability Zone A



Host 1



Host 2



Host 3

Availability Zone B



Host 4



Host 5



Host 6

Availability Zone C



Host 7



Host 8



Host 9

CustomerOrdersTable

Backup and restore

The only cloud database to provide on-demand and continuous backups



On-demand
backups for long-
term data archival
and compliance



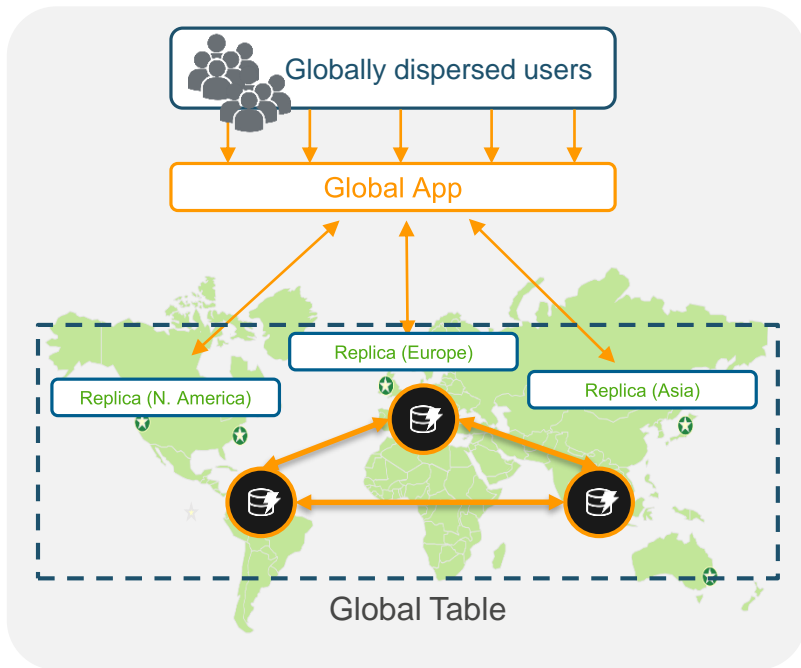
Point in time restore
for short term
retention and data
corruption protection
(35 days)



Point in time recovery with
restore times in a few hours
depending on table size

Global Tables

The first fully-managed, multi-master, multi-region database



Build high performance, globally distributed applications


Low latency reads & writes to locally available tables

Disaster proof with multi-region redundancy

Easy to setup and no application re-writes required

Capacity managed for you

Provisioned

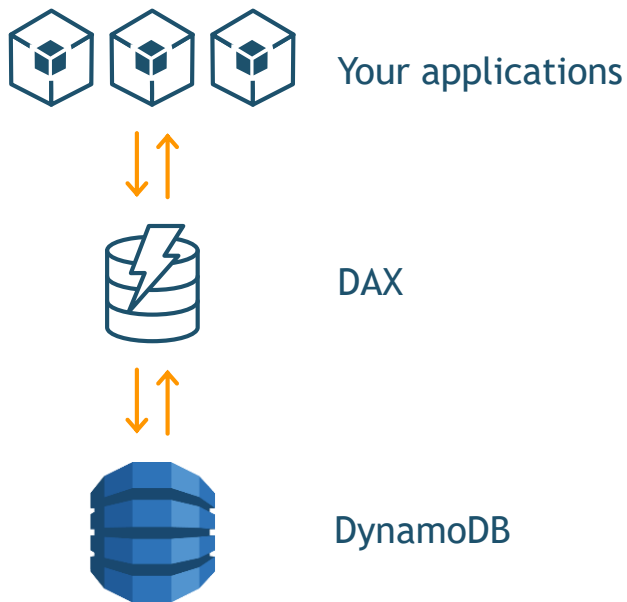


On-Demand



DynamoDB Accelerator (DAX)

High performance




Fully managed, highly available cache for DynamoDB

Even faster—microsecond latency

Scales to millions of requests per second

API compatible

Fully managed auto scaling



Automated
scaling policies

Scales up when
you need it


Scales down when
you don't

Scheduled
auto scaling

NoSQL vs. SQL for a new app: how to choose?

Want simplest possible DB management?


Want app to manage DB integrity?



Need joins, transactions, frequent table scans?

Want DB engine to manage DB integrity?

Team has SQL skills?



Task: Build a API based simple backend for Weather App

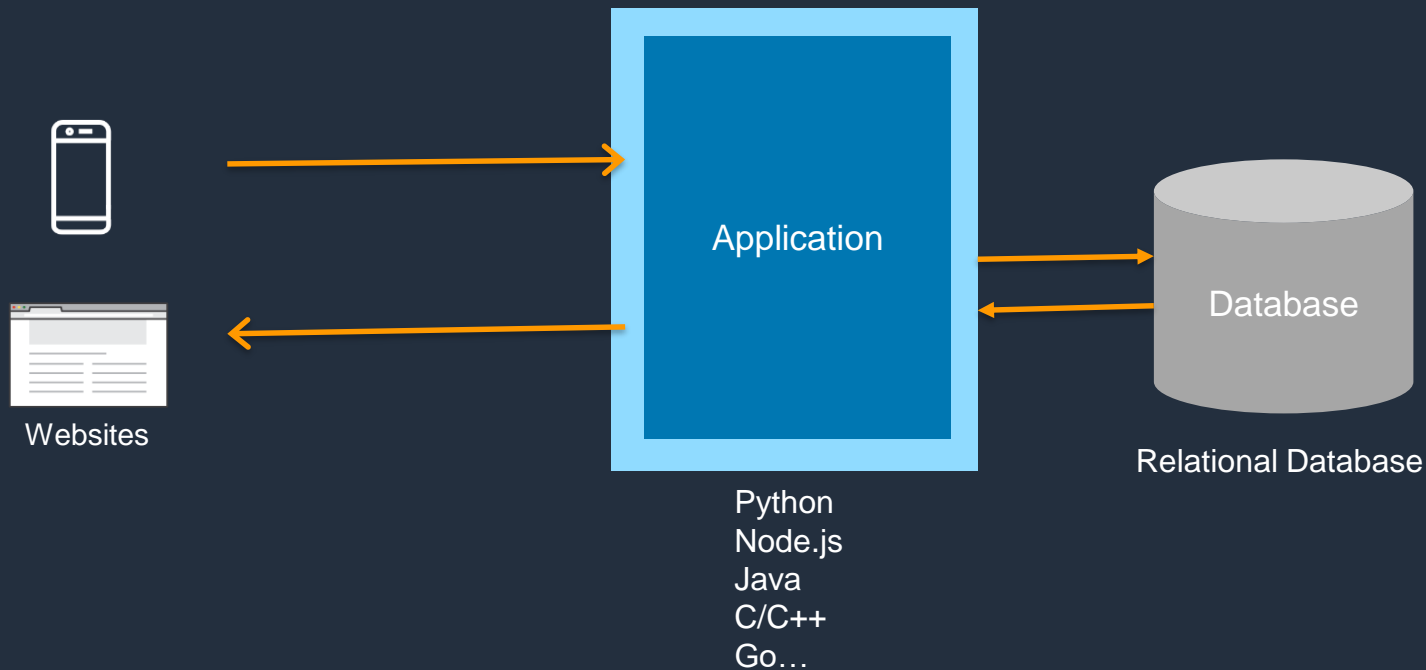
Daily Weather

Show Current

Pune ▼

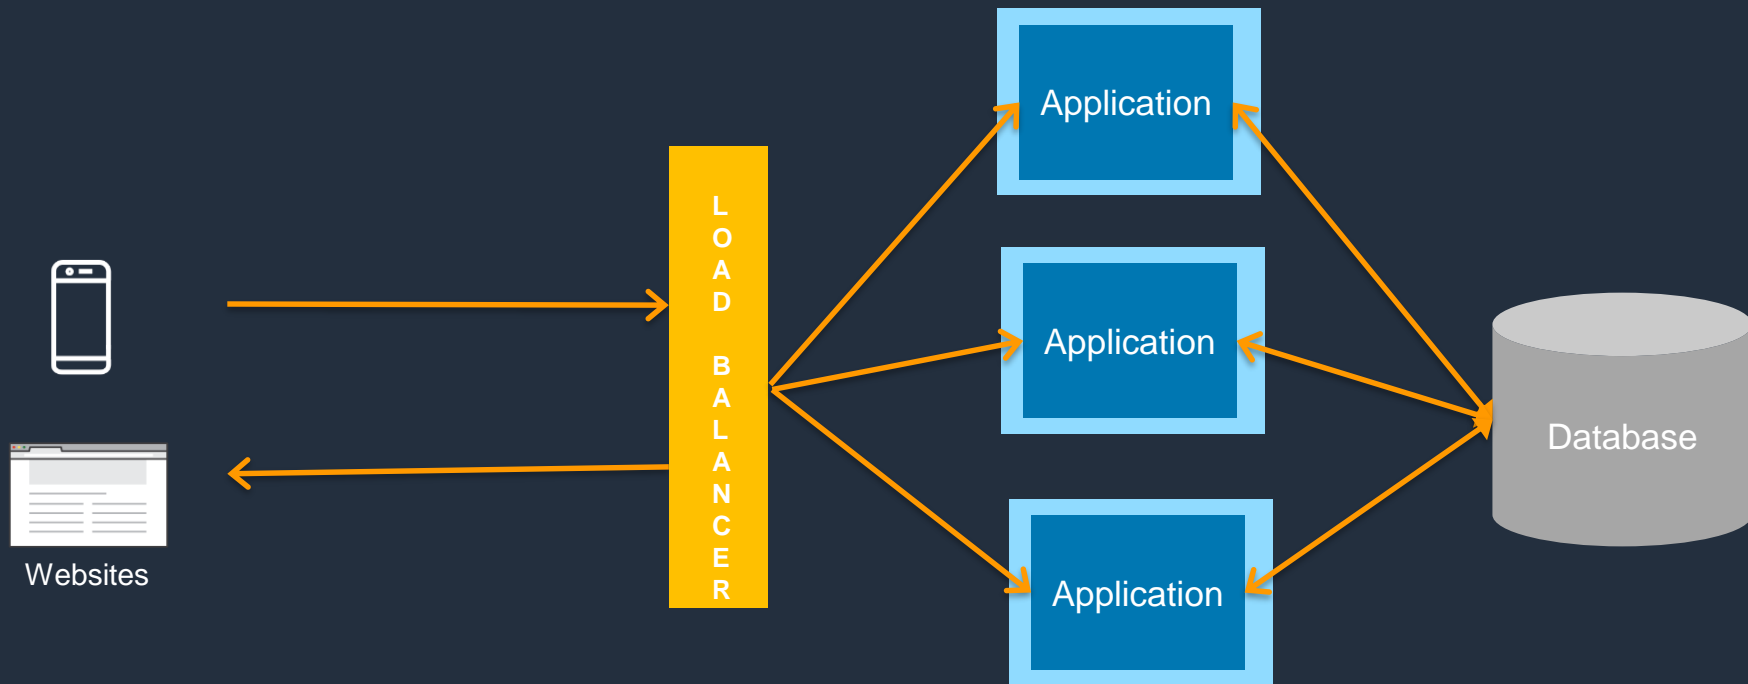
Value: 37 °C

Task: Build a API based simple backend for Weather App

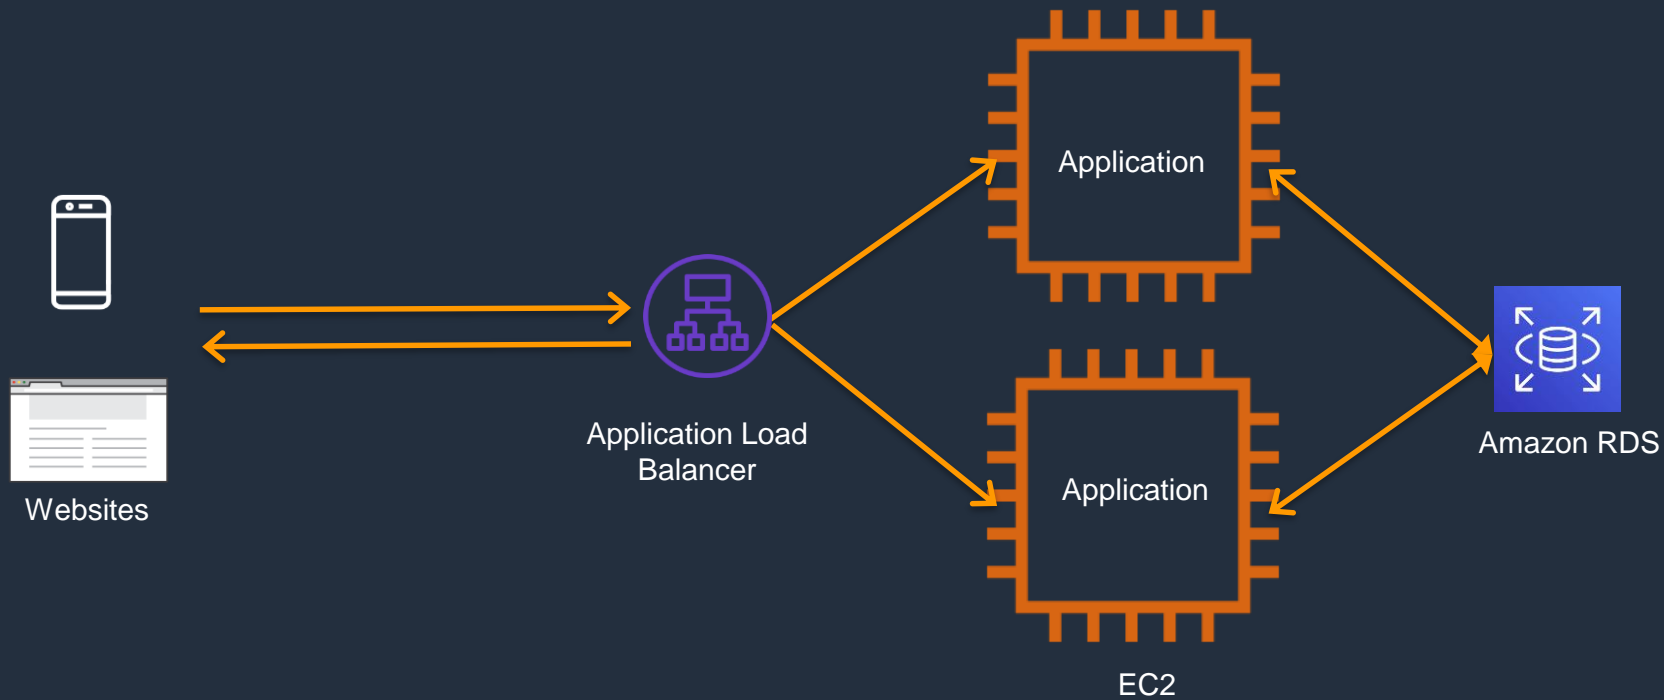


© 2018, Amazon Web Services, Inc. or its affiliates. All rights reserved.

How do you scale your application?



Let's host this on AWS using IaaS services (EC2)

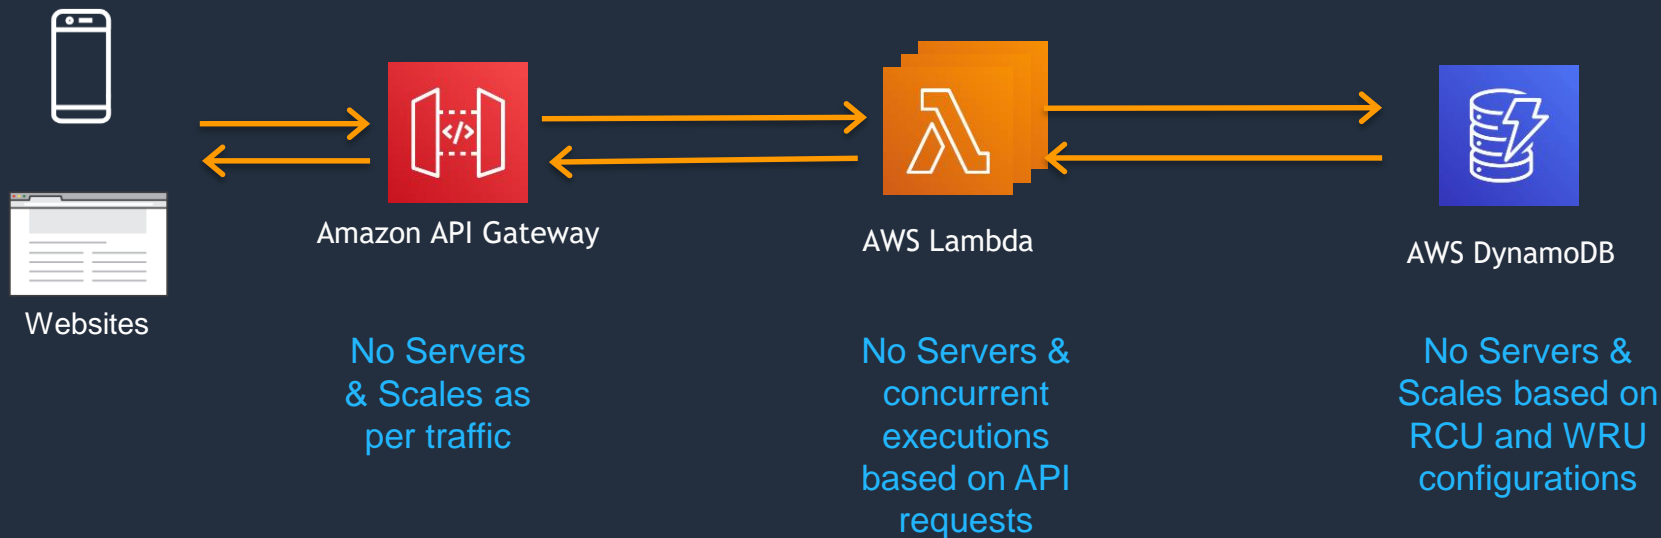


Do you see a problem with these architectures?

- Think about Scalability
- Think about Reliability
- Think about Performance
- Think about Operational efficiency
- **Think about Cost**

Solution: Let's host this on AWS using Serverless services

<http://d1osdxlszoytz4.cloudfront.net/>



Introducing Amazon ElastiCache

Fully-managed, Redis or Memcached compatible, low-latency, in-memory data store




Extreme Performance

In-memory data store and cache for sub-millisecond response times



Fully Managed

AWS manages all hardware and software setup, configuration, monitoring

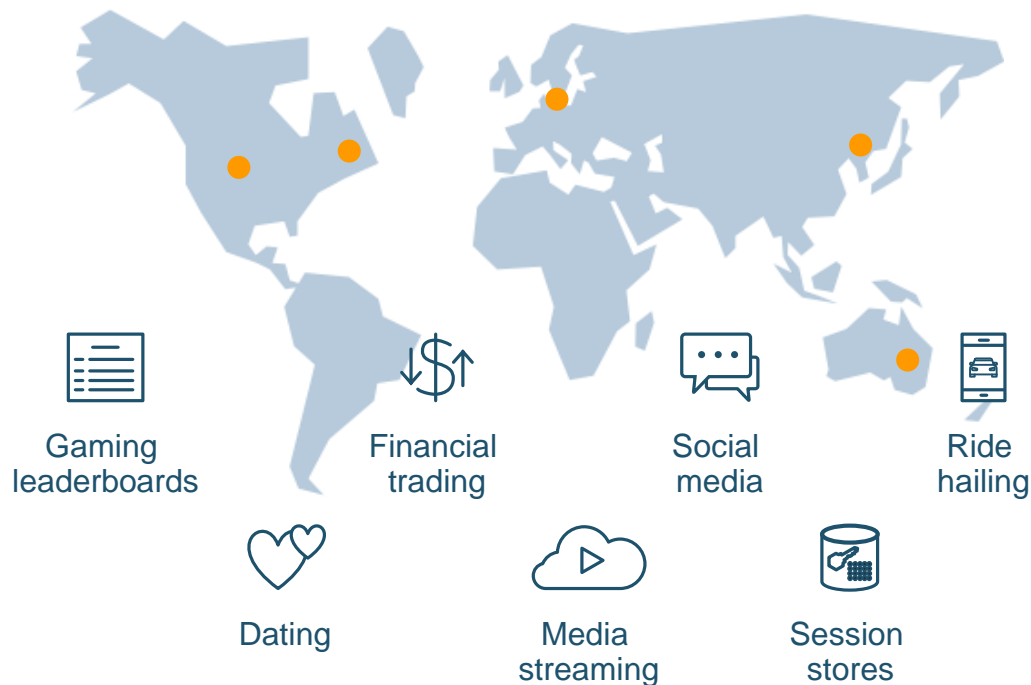


Easily Scalable

Read scaling with replicas. Write and memory scaling with sharding. Non disruptive scaling


μ s is the new *ms*

Internet-scale apps need low latency and high concurrency



Users	1M+
Data volume	TB-PB-EB
Locality	Global
Performance	Milliseconds to microseconds
Request Rate	Millions
Access	Mobile, IoT, Devices
Scale	Up-Out-In
Economics	Pay as you go
Developer access	Instant API access


Developers use various approaches to reduce latency



In-memory databases
and data grids



Specialized hardware
such
as multi-core processors,
GPUs, accelerators



Data reduction
approaches
such as sampling,
aggregation

Amazon ElastiCache

- In-memory cache in the cloud
- Improve latency and throughput for read-heavy workloads
- Supports open-source caching engines
 - Memcached
 - Redis
- Fully managed
- Multi-AZ

Examples

- Caching of MySQL database query results
- Caching of post-processing results
- Caching of user session and frequently accessed data



redis



ElastiCache Redis

#1 Key-Value Store*

Fast in-memory data store in the cloud. Use as a database, cache, message broker, queue

Highly Available & Reliable

Read replicas, multiple primaries, multi-AZ with automatic failover

Fully Managed & Hardened

AWS manages hardware, software, setup, configuration, monitoring, failure recovery, and backups

Easily Scalable

Cluster with up to 6.1 TiB of in-memory data

Read scaling with replicas

Write and memory scaling with sharding

Scale out or in

Secure & Compliant

VPC for cluster isolation, encryption at rest/transit, HIPAA compliance

*: <https://db-engines.com/en/ranking>




ElastiCache Memcached



Fully Managed
Memcached

Fast in-memory data store in the cloud. Use as a cache to reduce latency and improve throughput




Secure &
Hardened

VPC for cluster isolation



Easily
Scalable

Sharding to scale in-memory cache with up to 20 nodes and 8.14 TiB per cluster



Fully managed graph database

Supports open graph APIs

Scalable

ACID compliant

Multi-AZ

Amazon Neptune

Fully managed graph database for highly connected data

Open



Supports Apache TinkerPop™ & W3C RDF graph models

Fast & Scalable




Store billions of relationships; query with millisecond latency

Reliable



6 replicas of your data across 3 AZs with full backup and restore

Easy





Build powerful queries easily with Gremlin and SPARQL
+
GRAPHQL with AppSync



Use cases for highly connected data

- Social networking
- Recommendations
- Knowledge graphs
- Fraud detection
- Life sciences
- Network and IT operations





Fully managed time series database

1,000x faster at 1/10th the cost

Built-in analytics

Serverless

Building with time-series data is challenging

Relational databases



Unnatural for
time-series
data



Inefficient
time-series data
processing




Rigid schema
inflexible for fast
moving time-series
data

Existing time-series databases



Difficult to
scale



Difficult to
maintain high
availability



Limited data
lifecycle
management

Amazon Timestream

(Preview)

**1,000x faster at 1/10th
the cost of relational
databases**



Collect fast moving time-series data from multiple sources at the rate of millions of inserts per second

**Trillions of daily
events**



Capable of processing trillions of events daily; the adaptive query processing engine maintains steady, predictable performance

**Analytics
optimized for time
series data**




Built-in analytics for interpolation, smoothing, and approximation to identify trends, patterns, and anomalies

Serverless



No servers to manage; time-consuming tasks such as hardware provisioning, software patching, setup, & configuration done for you



Fully managed ledger database

Immutable and transparent

Cryptographically verifiable

Scalable

Serverless

Amazon Quantum Ledger Database (QLDB)

Fully managed ledger database

Track and verify history of all changes made to your application's data

Immutable




Maintains a sequenced record of all changes to your data, which cannot be deleted or modified; you have the ability to query and analyze the full history

Cryptographically verifiable



Uses cryptography to generate a secure output file of your data's history

Highly scalable



Executes 2–3X as many transactions as ledgers in common blockchain frameworks

Easy to use



Easy to use, letting you use familiar database capabilities like SQL APIs for querying the data

Common customer use cases



Banking & Finance

Keeping track of transactions,
trades and accounts



E-Commerce

Where's my stuff?



Transport & Logistics

Tracking transportation
of goods



HR & Payroll

Tracking changes to an
individual's profile



Manufacturing


Recording components used
in manufacturing



Government

Tracking vehicle
title history

[Now GA]



Fully managed document database

MongoDB-compatible

Multi-AZ

Performance at scale

Why use a document database?

The JSON document model maps naturally to application data



Each document can have a different data structure and is independent of other documents




Index on any key in a document, and run ad hoc and aggregation queries across your data set



Use cases for document databases



Content
Management




Mobile




Personalization



Catalog



Retail and
Marketing



User profiles

Use case: Profile Management

users table

id	username	first_name	last_name
181276	sue1942	Susan	Benoit

tankfight_users table

id	hi_score	global_rank
181276	3185400	5139

```
{  
  id: 181276,  
  username: 'sue1942',  
  name: {first: 'Susan',  
         last: 'Benoit'},  
}
```



```
{  
  id: 181276,  
  username: 'sue1942',  
  name: {first: 'Susan',  
         last: 'Benoit'},  
  tankfight: {  
    hi_score: 3185400,  
    global_rank: 5139  
  }  
}
```

It's all about choice

Performance-oriented
Cost-oriented

Common data categories and use cases



Relational

Referential integrity, ACID transactions, schema-on-write



Key-value

High throughput, low-latency reads and writes, endless scale



Document

Store JSON documents with quick access, query on any attribute



In-memory

Query by key with microsecond latency



Graph

Quickly and easily create and navigate relationships between data



Time-series

Collect, store, and process data sequenced by time



Ledger

Complete, immutable, and verifiable history of all changes to application data



Warehouse

High performance querying on large volumes of data



RDS



DynamoDB



DocumentDB



ElastiCache

Redis Memcached



Neptune



Timestream



QLDB



Redshift

Thank you!

Chetan Agrawal – agrcheta@amazon.com